

Nischal Chandur

858-241-6448 | chandur.nischal2@gmail.com

nchandur.github.io/portfolio | linkedin.com/in/nchandur | github.com/nchandur

SUMMARY

Machine Learning Engineer with 2+ years of experience designing and deploying production-grade ML systems, with specialization in LLMs, RAG architectures, and geospatial data science. Skilled in transforming cutting-edge NLP and ML research into real-world, trustworthy products. Proven ability to lead end-to-end pipelines, from model development and cloud deployment to cross-functional collaboration with domain experts and business stakeholders. Strong focus on interpretability, ethical AI, and secure, scalable systems.

EXPERIENCE

Data Science Graduate Intern, Ecolab – Naperville, IL, USA

06/2024 – 08/2024

- Built a RAG assistant using Azure AI Studio to enable store managers to query KPIs from structured uploads (e.g., Excel).
- Designed an LLM-based support tool for field agents using embeddings of manuals and incident logs for contextual troubleshooting.
- Presented solution demos and architectural decisions to executives; worked with engineers to refine requirements.
- Created secure NLP interfaces for natural language interaction over mixed-format data.

Machine Learning Engineer, Reworked.ai – Miami, FL, USA

04/2024 – 05/2024

- Integrated geospatial features into the ML pipeline, enhancing the targeting strategy and improving customer segmentation
- Developed a Bayesian ensemble model to predict solar panel adoption likelihood, contributing to smarter customer acquisition and outreach strategies.
- Communicated value propositions to leadership through strategic insights and data-driven recommendations.

Data Scientist, Latlong – Bengaluru, KA, India

09/2022 – 06/2023

- Led development of a multilingual OCR-NLP pipeline for structured data extraction from public government documents, enabling regional demographic analytics for clients.
- Engineered geospatial risk models for NBFCs to assess borrower reliability by fusing behavioral and location-derived features into credit scoring.
- Delivered spatial analytics to a major automotive client, correlating product performance and sales patterns with demographic and location data.
- Presented findings and deployment options directly to product heads and senior business stakeholders.

PROJECTS

Lorekeeper – RAG-based Q&A System, University of Maryland

08/2024 – 12/2024

- Developed a production-style Retrieval-Augmented Generation (RAG) system using Llama 3.2:1b, LangChain, and FAISS, enabling interpretable, real-time question answering over book corpora.
- Designed chunking, embedding, and search modules with modularity and scalability in mind.
- Implemented a front-end with source traceability and explanation features to enhance trust and transparency in LLM outputs.
github.com/nchandur/lorekeeper

NBA Game Outcome Predictor & Analytics Dashboard, University of Maryland

08/2023 – 12/2024

- Built an end-to-end ML pipeline for predicting NBA game outcomes using ensemble models trained on 40+ years of historical data.
- Led team efforts in model design, REST API integration, and dashboard development using Flask.
- Focused on explainability and usability by surfacing key player metrics, model confidence intervals, and interactive visualizations.
github.com/nchandur/NBA-prediction-model

EDUCATION

University of Maryland, College Park, MD, United States

08/2023 – 05/2025

Master of Science in Data Science

GPA: 3.9

Coursework: Natural Language Processing, Computer Vision, Big Data Systems, Algorithms for Data Science, Cloud Computing

SKILLS

Python, R, C/C++, Go, JavaScript, Bash, HTML/CSS, LangChain, Hugging Face, ChromaDB, FAISS, Sentence Transformers, Pinecone, Weaviate, Neo4j, OCR, Prompt Engineering, PyTorch, TensorFlow, Scikit-learn, Keras, Bayesian Inference, Ensemble Learning, Docker, Azure AI Studio, Amazon Web Services, Streamlit, Flask, REST APIs, CI/CD, Git, Spark, Hadoop, PostgreSQL, MongoDB, Redis, Dask, NumPy, Pandas, Matplotlib, Seaborn, Plotly, Data Cleaning, Feature Engineering, Exploratory Data Analysis, Statistical Testing, Time Series Analysis, Airflow, Kafka, Snowflake, BigQuery, Data Modeling, ETL/ELT, Data Warehousing, Parquet